

Homework 2

Benny Chen

September 20, 2022

Question 1

Suppose A and B are word arrays. The following C loop increments elements in A by 4 and saves the results into B.

```
for (i = 0; i < 100; i += 1)
    B[i] = A[i] + 4;
```

The following table shows the mapping between variables and registers.

Register	s1	s2	s3
Variable/value	i	Address of A	Address of B

We will study two implementations in RISC-V.

- a) The first implementation is based on the array copy code we discussed in lecture. We just need to revise it slightly. What changes do we need? How many instructions will be executed for the loop? Note that we do not need to jump to the condition test before the first iteration because we are sure the condition is true at the beginning.

Part A answer:

We are already given the bulk of the code from the lectures. The only change that we would need to make is to add 4 to A[i] after we load the value to t1. There would be a total of 802 executed instructions.

```
addi    s4, x0, 100
addi    s1, x0, 0
beq     x0, x0, test # we know s1 < s4

loop:

slli    t0, s1, 2 # t0 = i * 4
add     t2, t0, s2 # compute addr of A[i]
lw      t1, 0(t2)
addi    t1, t1, 4 # add 4 to A[i]
add     t3, t0, s3 # compute addr of B[i]
```

```

        sw      t1, 0(t3)
        addi    s1, s1, 1

test:
        bne     s1, s4, loop # 7 instructions in the loop

```

- b) Loop unrolling is an optimization technique to improve the performance of programs. In the second implementation, we unroll the loop and process four array elements in A in each iteration. The unrolled loop in C is shown below. Translate the loop to RISC-V instructions. Try to minimize the number of instructions that are executed. Explain your code. How many instructions will be executed for the new loop?

```

for (i = 0; i < 100; i += 4) {
    B[i] = A[i] + 4;
    B[i+1] = A[i+1] + 4;
    B[i+2] = A[i+2] + 4;
    B[i+3] = A[i+3] + 4;
}

```

Part B answer:

The problem is similar to part A so we can start with a baseline from Part A. The change for Part B is that we are adding 4 to the index (i) and loading 4 values from A[i,i+1,i+2,i+3] to B[i,i+1,i+2,i+3]. Since we know how to find the address of A[i] and B[i] we can just add 4 to the address to increment the index by 1. There would be a total of 427 executed instructions.

```

        addi    s4, x0, 100 #100
        addi    s1, x0, 0 #i = 0
        beq     x0, x0, test # we know s1 < s4

loop:
        #Address calculation
        slli    t0, s1, 2 # t0 = i * 4

        # B[i] = A[i] + 4
        add     t2, t0, s2 # compute addr of A[i]
        lw      t1, 0(t2)
        addi    t1, t1, 4 # add 4 to A[i]
        add     t3, t0, s3 # compute addr of B[i]
        sw      t1, 0(t3) # load A[i] + 4 to B[i]
        # We know A[i] and B[i] address so we just need
        # to load address + 4 to have A and B [i+1]
        # B[i + 1] = A[i + 1] + 4
        lw      t1, 4(t2) # load A[i + 1]
        addi    t1, t1, 4 # add 4 to A[i + 1]
        sw      t1, 4(t3) # load A[i + 1] + 4 to B[i + 1]

```

```

# B[i + 2] = A[i + 2] + 4
lw      t1, 8(t2) # load A[i + 2]
addi    t1, t1, 4 # add 4 to A[i + 2]
sw      t1, 8(t3) # load A[i + 2] + 4 to B[i + 2]

# B[i + 3] = A[i + 3] + 4
lw      t1, 12(t2) # load A[i + 3]
addi    t1, t1, 4 # add 4 to A[i + 3]
sw      t1, 12(t3) # load A[i + 3] + 4 to B[i + 3]

addi    s1, s1, 4 # i+=4

test:
      bne      s1, s4, loop # 7 instructions in the loop

```

Question 2

A two-dimensional array in C (and some other languages) can be considered as an array of one-dimensional array. For example, the following define T as an 16x8 array in C.

```
int T[16][8];
```

The two-dimensional array can be considered as an array of 16 elements, each of which is a one-dimensional array of 8 integers/words. In total there are 128 words. The words are stored in memory in the following order:

```

T[0][0], T[0][1], ..., T[0][6], T[0][7],
T[1][0], T[1][1], ..., T[1][6], T[1][7],
...
T[14][0], T[14][1], ..., T[14][6], T[14][7],
T[15][0], T[15][1], ..., T[15][6], T[15][7]

```

Row 0, consisting of $T[0][0]$, $T[0][1]$, ..., and $T[0][7]$, goes first. Row i is stored right after row $i - 1$, for $i = 1, 2, \dots, 15$. For example, $T[1][0]$ is stored right after $T[0][7]$. If $T[0][0]$ is located at address 1000, $T[0][7]$ is located at address $1028 = 1000 + 7 * 4$. And $T[1][0]$ is located at address 1032. Similarly, we can calculate that $T[2][0]$ is located at 1064, $T[3][0]$ is located at 1096, and so on.

Translate the following C code to RISC-V instructions. Assume T's address is already in s9. As a practice of accessing two-dimensional arrays, do not use pointers. Explain your code, especially how you implement the loops and how you calculate $T[i][j]$'s address.

```

for (i = 0; i < 16; i += 1)
    for (j = 0; j < 8; j += 1)
        T[i][j] = 256 * i + j;

```

Answer:

The original C code basically has a double for loop that in the first loop, iterates 16 times and in the second loop 8 times, all the while setting the value $T[i][j]$ to $256 * i + j$. The RISC-V code works in a similar way. We first defined the values of i and the constants like 16 and 8. We then jump to the end where the test case is and checks if $i \leq 16$ while setting j to 0. If the test case is true, we jump to the inner loop where we perform the operation to put in the value for $256 * i + j$. We then find the address of $T[i][j]$ by first computing the address of $T[i]$ then adding the address for j into it. We then store the value into the address then increment j and repeat the loop until all statements are true and executed.

```
# s9 = T's address

addi    s0,s0,0 #i = 0
addi    s2,s2,16 #16
addi    s3,s3,8 #8
beq     x0,x0,test

action:  slli    t0,s0,256 # t0 = i * 256
        add     t0,t0,s1 # t0 = t0 + j

        slli    t1,s0,2 # t1 = i * 4
        slli    t2,s1,2 # t2 = j * 4
        add     t3,t1,s9 # Compute address of T[i]
        add     t3,t3,t2 # Compute address of T[i][j]
        sw     t0,0(t3) # T[i][j] = t0

test2:   addi    s1,s1,1 #j++
        blt     s1,s3,action #j < 8
        addi    s0,s0,1 #i++

test:    addi    s1,x0,0 #j = 0
        blt     s0,s2,action # i < 16
```

Question 3

Encoding. For each RISC-V instruction, find out its encoding format, *the bits* in each field, and the machine code as 8 hexadecimal digits. An example is shown below. Pay attention to the number of bits in each field.

```
or      s1, s2, s3
slli    t1, t2, 16
xori    x1, x1, -1
lw      x2, -100(x3)
```

Answer:

CORE INSTRUCTION FORMATS

	31	27	26	25	24	20	19	15	14	12	11	7	6	0
R	funct7				rs2		rs1		funct3		rd		Opcode	
I	imm[11:0]					rs1		funct3		rd		Opcode		
S	imm[11:5]				rs2		rs1		funct3		imm[4:0]		opcode	
SB	imm[12 10:5]				rs2		rs1		funct3		imm[4:1 11]		opcode	
U	imm[31:12]										rd		opcode	
UJ	imm[20 10:1 11 19:12]										rd		opcode	

REGISTER	NAME	USE	SAVER
x0	zero	The constant value 0	N.A.
x1	ra	Return address	Caller
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	--
x4	tp	Thread pointer	--
x5-x7	t0-t2	Temporaries	Caller
x8	s0/fp	Saved register/Frame pointer	Callee
x9	s1	Saved register	Callee
x10-x11	a0-a1	Function arguments/Return values	Caller
x12-x17	a2-a7	Function arguments	Caller
x18-x27	s2-s11	Saved registers	Callee
x28-x31	t3-t6	Temporaries	Caller
f0-f7	ft0-ft7	FP Temporaries	Caller
f8-f9	fs0-fs1	FP Saved registers	Callee
f10-f11	fa0-fa1	FP Function arguments/Return values	Caller
f12-f17	fa2-fa7	FP Function arguments	Caller
f18-f27	fs2-fs11	FP Saved registers	Callee
f28-f31	ft8-ft11	R[rd] = R[rs1] + R[rs2]	Caller

Instruction: or s1, s2, s3

Register numbers: or x9, x18, x19

Type: R-Type

opcode: 0110011

rd: 01001

funct3: 110

rs1: 10010

rs2: 10011

funct7: 0000000

Machine code in bits: 0000 0001 0011 1001 0110 0100 1011 0011

Machine code in hex: 0x013964B3

or R 0110011 110 0000000 33/6/00

funct7	rs2	rs1	funct3	rd	Opcode
0000000	10011	10010	110	01001	0110011

Instruction: slli t1, t2, 16

Register numbers: slli x6,x7,16

Type: I-Type

opcode: 0010011

rd: 00110

funct3: 001

rs1: 00111

shift: 10000

Immediate: 0000000

Machine code in bits: 0000 0001 0000 0011 1001 0011 0001 0011

Machine code in hex: 0x1039313

```

slli          I      0010011      001      0000000      13/1/00

```

imm[11:0]	rs1	funct3	rd	Opcode
000000010000	00111	001	00110	0010011

Instruction: xori x1, x1, -1

Register numbers: xori x1,x1,-1

Type: I-Type

opcode: 0010011

rd: 00001

funct3: 100

rs1: 00001

Immediate: 111111111111

Machine code in bits: 1111 1111 1111 0000 1100 0000 1001 0011

Machine code in hex: 0xFFFF0C093

```

xori          I      0010011      100      111111111111      13/4

```

imm[11:0]	rs1	funct3	rd	Opcode
111111111111	00001	100	00001	0010011

Instruction: lw x2, -100(x3)

Register numbers: lw x2, -100(x3)

Type: I-Type

opcode: 0000011

rd: 00010

funct3: 010

rs1: 00011

Immediate: 111110011100

Machine code in bits: 1111 1001 1100 0001 1010 0001 0000 0011

Machine code in hex: 0xF9C1A103

lw I 0000011 010 03/2

imm[11:0]	rs1	funct3	rd	Opcode
111110011100	00011	010	00010	0000011

Question 4

Decoding. Each 8-digit hexadecimal number in the following table represents a RISC-V instruction. For each machine code, find its encoding format, bits in each field, and then decode it into a RISC-V instruction. The steps are the reverse of those in Problem 3. Use register numbers (like x0 instead of zero). Any immediate or displacement (offset) should be in decimal.

	Machine Code
a	0xfeaca823
b	0x04020713
c	0x00557bb3
d	0x414fdf13

Answer:

Machine code: 0xFEACA823

sw S 0100011 010 23/2

Code in bits: 1111111 01010 11001 010 10000 0100011

Type: S-Type

Opcode: 0100011 (Type S-Type)

Immediate1: 10000 (Decimal: 16)

funct3: 010 (Instruction sw)

rs1: 11001 (Decimal 25)

rs2: 01010 (Decimal 10)

Immediate2: 1111111 (Negative)

Translation: sw x10, -16(x25)

Machine code: 0x04020713

addi	I	0010011	000		13/0
------	---	---------	-----	--	------

Code in bits: 000001000000 00100 000 01110 0010011

Type: I-Type

Opcode: 0010011 (Type I-Type)

rd: 01110 (Decimal 14)

funct3: 000 (Instruction addi)

rs1: 00100 (Decimal 4)

Immediate: 000001000000 (Decimal 64)

Translation: addi x14, x4, 64

Machine code: 0x00557BB3

and	R	0110011	111	0000000	33/7/00
-----	---	---------	-----	---------	---------

Code in bits: 00000000 00101 01010 111 10111 0110011

Type: R-Type

Opcode: 0110011 (Type R-Type)

rd: 10111 (Decimal 23)

funct3: 111 (Instruction and)

rs1: 01010 (Decimal 10)

rs2: 00101 (Decimal 5)

funct7: 0000000 (Decimal 0)

Translation: and x23, x10, x5

Machine code: 0x414FDF13

srai	I	0010011	101	0100000	13/5/20
------	---	---------	-----	---------	---------

Code in bits: 0100000010100 11111 101 11110 0010011

Type: I-Type

Opcode: 0010011 (Type I-Type)

rd: 11110 (Decimal 30)

funct3: 101 (Instruction srai)

rs1: 11111 (Decimal 31)

shift: 010100 (Decimal 20)

Immediate: 0100000

Translation: srai x30, x31, 20