

# Final Exam Topics

## Numbers, bits, bytes, and ASCII characters

- Binary numbers, two's complement numbers, hexadecimal numbers.
- Conversion between different number systems, including decimal numbers, binary, hexadecimal.
- Sign extension.
- Addition/subtraction of binary/two's complement numbers. Negate a two's complement number.
- Bitwise logical operations.
- ASCII characters.

## RISC-V ISA

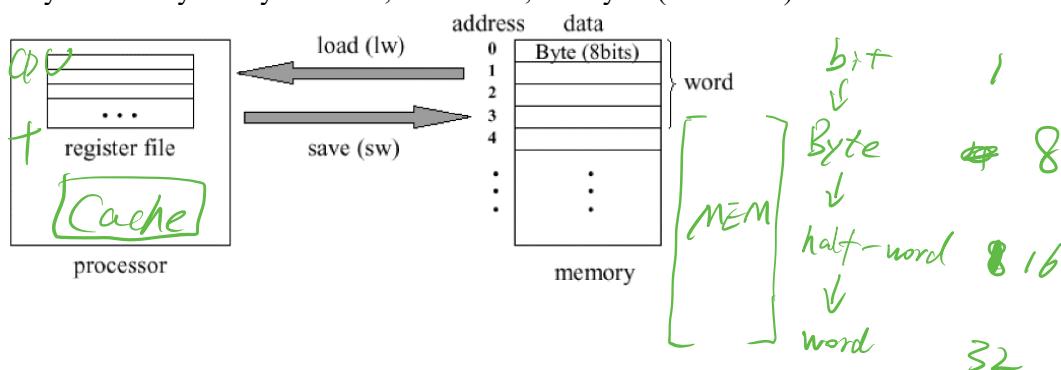
- RISC-V instruction sets. Arithmetic, logical, memory, control flow instructions.

Instruction class	RISC-V examples
Arithmetic	add, sub, addi...
Data transfer	lw, sw, lb, lbu, sb lh, lhu, lui...
Logical	and, or, xor, andi, ori, sll, sra...
Cond. Branch	beq, bne...
Jump	Jal, jalr...

## Data

*CPU → Cache → MEM*

- Accessing data in registers and memory.
- Array in memory. Array of words, half-words, and bytes (characters).



*[Cache] block = 1 word  
[ ] 2 words  
N words*



## Control flow

- Program counter.
- RISC-V support for if-then-else, loops, functions.
- Compare values with beq, bne, etc.
- RISC-V calling convention. Caller-saved / callee-saved registers. Passing parameters to / returning value from functions.
- Stack. Push/pop. Save/restore registers on stack. Allocate storage space on stack.

## Instruction Encoding

- Six instruction formats (R, I, S, SB, U, UJ types).
- Fields in different instruction formats.
- Describe the placement of immediate bits in machine code, for different types of formats.
- Describe how assembler place bits in immediate in machine code and how the processor construct.
- immediate when executing instructions.
- Given enough information, encode RISC-V instruction with 32 bits and read encoded instructions.

## Assembly code

- Translate pseudocode or simple C code to RISC-V assembly code.
- Understand/debug RISC-V code.

## Arithmetic for Computers

- operations on integers
  - addition and subtraction;
  - multiplication and division; multiplication/division algorithm and hardware;
  - RISC-V Multiplication Instructions
- floating-point and real numbers
  - Binary to decimal; Decimal to binary
  - Normalized numbers.
  - Encode Floating Point Numbers with Bits
  - IEEE Std 754-1985: Single precision; double precision
  - Denormal Numbers; meaning.
  - FP Support in RISC-V

## Introduction to Digital Design

- use Boolean algebra to describe circuits
  - Boolean operations: AND, OR, NAND, NOR, XOR, NOT
  - use of truth tables to describe functionality

- o Boolean algebra
  - o deriving Boolean expressions from truth tables
- decoder; multiplexor
- 1-bit ALU
  - o use Nx1 mux to select desired operation
  - o operations:
    - add, sub, AND, OR, NOR, etc. with individual gates
    - addition with full adder
    - subtraction with full adder and two's complement trick
- 32-bit ALU
- Describe the general structure of a state machine.
- Given a design of digital circuit, describe its behavior/function.
- Analyze the timing of digital circuit. Sequential Circuits
- Clock cycle, clock rate
- Bistable element, Memory Elements, Flip-Flops
- Triggering: falling edge and rising edge → structure hazard
- Registers: clock, input, output, control
  - o Shift Register
- Register Files

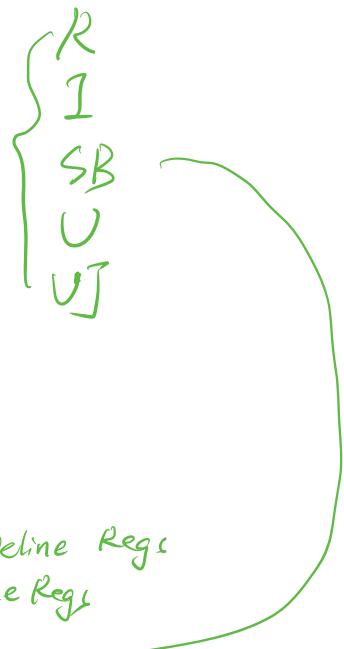
## Performance

- CPU Execution Time and Throughput; Relative Performance
- CPU Time = CPU Clock Cycles \* Clock Cycle Time = CPU Clock Cycles/Clock rate
- CPU Time = Instruction Count \* CPI \* Clock \* Cycle Time
  - o Instruction Count and CPI
  - o average CPI → Exam 2
- How to improve the performance of a computer/CPU? Amdahl's Law Exam 2

## Datapath Design

Fig 4.21

- general concepts
  - o tasks of "executing an instruction" into stages
  - o create the entire datapath for R/I type instruction.
- typical stages
  - o instruction fetch (IF) - all
  - o instruction decode (ID) / register fetch - all
  - o execution (EX/EXE) - varies based on different instruction
  - o memory access (MEM) - load and store only
  - o register write (WB) - ALU and load only
- standard components
  - o register file
  - o PC
  - o "add 4" unit, i.e., PC+4
  - o sign-extension
  - o instruction memory
  - o data memory
  - o ALU and its functions; ALU Control
  - o multiplexors
- implementation choices
  - o single-cycle: perform all tasks in a single clock cycle no pipeline Regs
  - o multi-cycle/pipeline: use a clock cycle for each stage Pipeline Regs



## Control Implementation

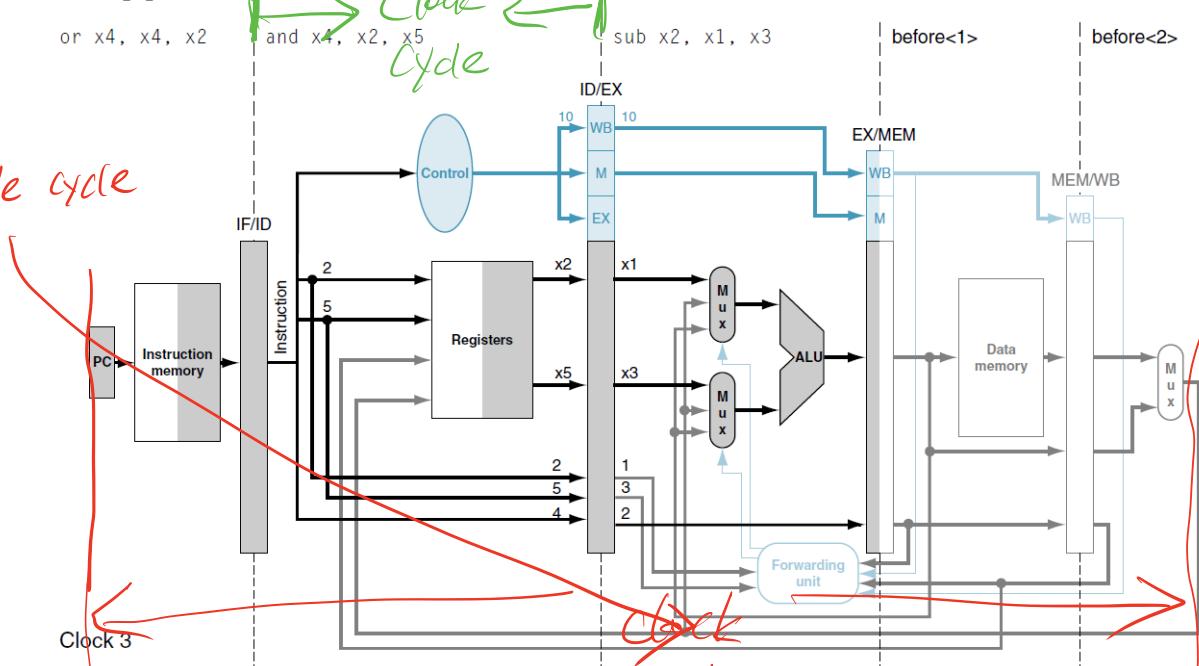
- single cycle control unit
  - o used for directing operations of datapath blocks

- generated based on opcode, function code (e.g., func 3 and func 7)
- figure out what control signals need to be activated for different R/I/S-type instructions.

## Pipelining

- basic concepts
- simple 5-stage pipeline
- pipeline registers
- pipeline control units

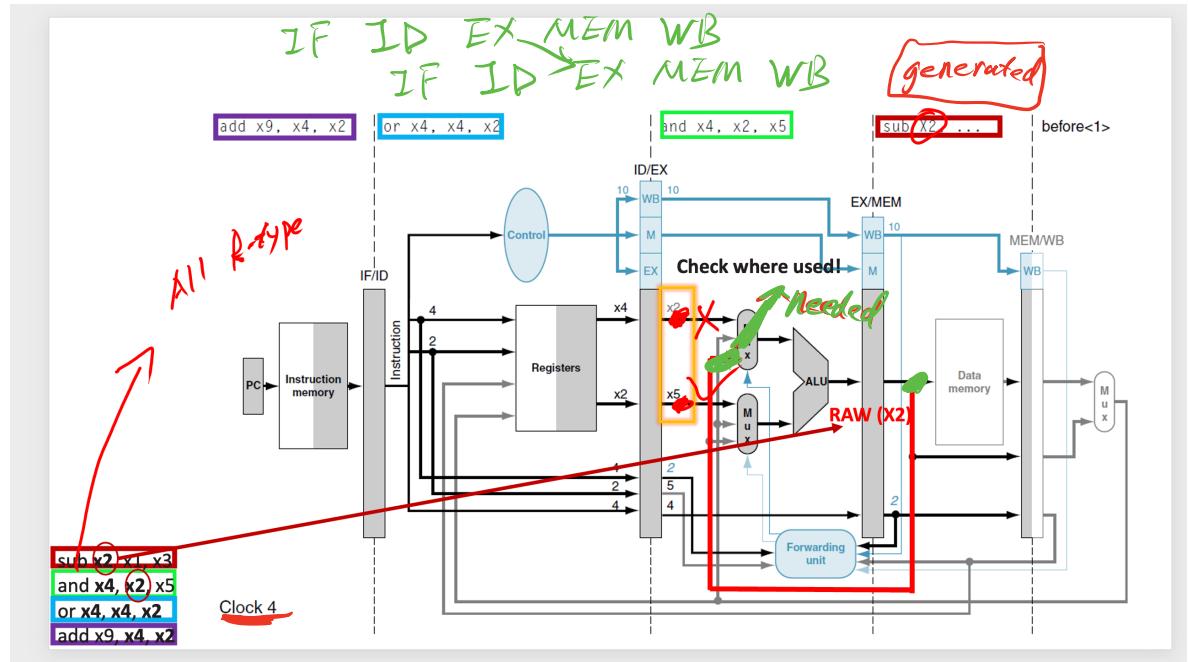
or  $x_4, x_4, x_2$



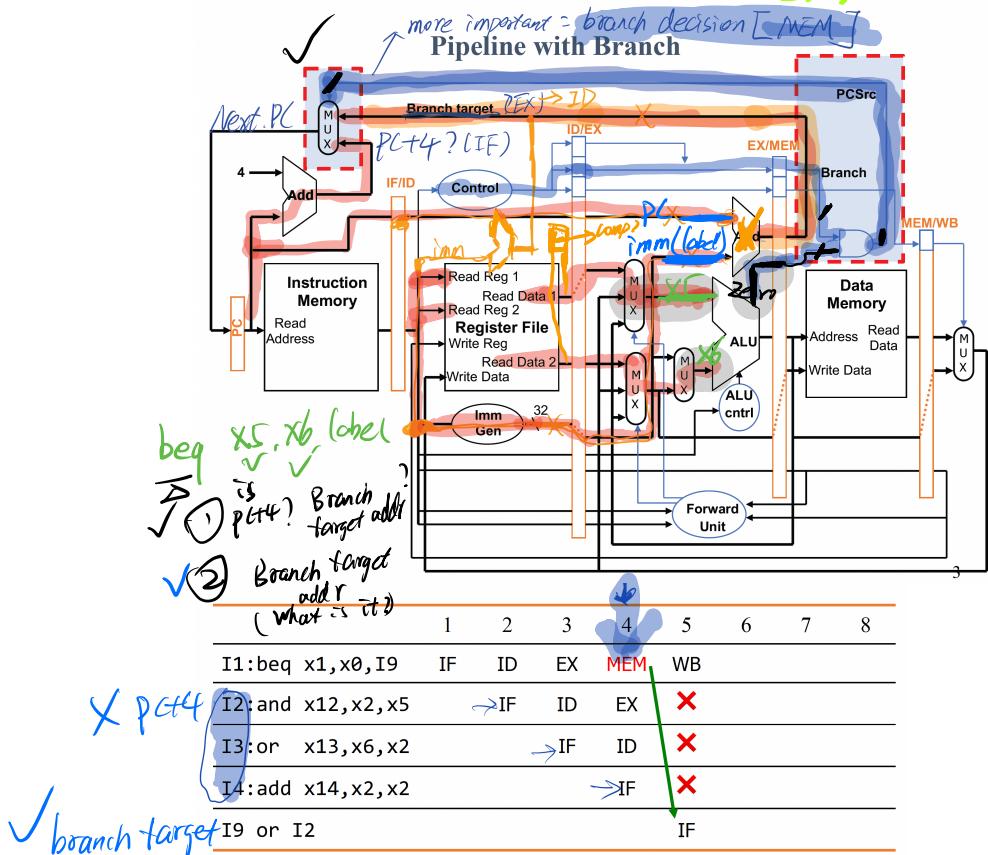
## Pipelining Hazard

- Structural hazard
  - Add more resources (IM, DM)
  - Share Resource (Reg)
- Data Hazard
  - RAW, Load-use, Mem-to-Mem copy
  - Forwarding
  - Stalling

F S



- Control Hazard
  - Stall (impacts CPI); Two “Types” of Stalls



- Move decision point as early in the pipeline as possible, thereby reducing the number of stall cycles
  - Cost of resolving branch in ID
- Branch prediction

- Static branch prediction
  - Predict not taken
  - Predict not taken works well for branching at “top of the loop”
    - But such loops have “jumps” at the bottom of the loop to return to the top of the loop and incur stalls
  - Predict not taken does NOT work well for branching at “bottom of the loop”
    - The branch is taken most times
    - Fall through instruction is always flushed
  - Predict taken

**Observation:** We can calculate the branch target address early.

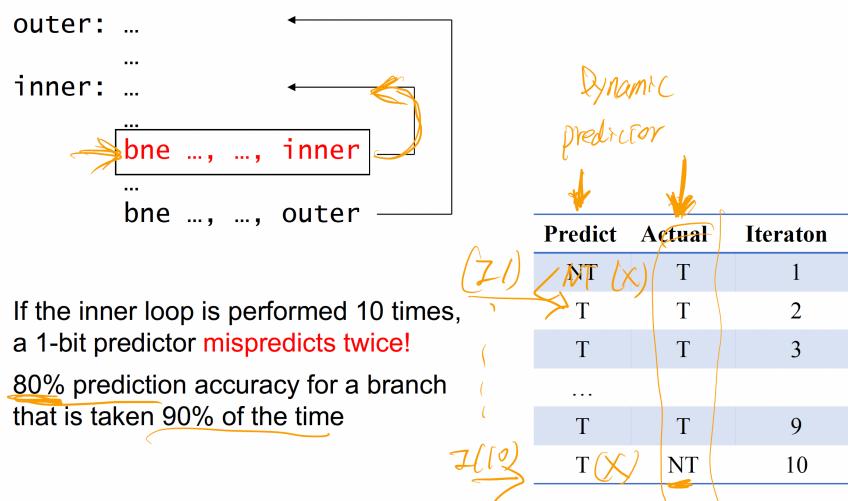
- It is easier to calculate branch target = PC + offset, because PC and offset are available early
- It is harder to read registers (data hazards!) and compare

**Predict taken:** predict branches will be taken

Once the branch target address is known, go to the branch target before the branch decision is made

- Need to know the branch target address. So predict taken *always* incurs (at least) one stall cycle
- Works for deeper pipelines, where the outcome of comparison is available in much later stages

- Dynamic branch prediction



- Re-order instructions to avoid hazards

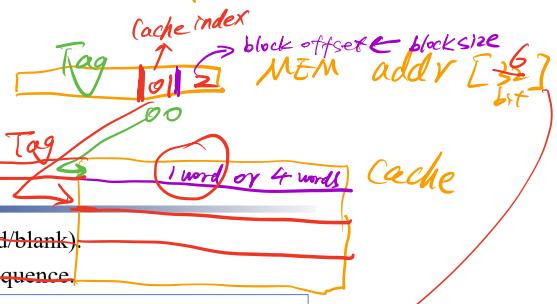
## Memory/Cache

- CPU/memory performance gap
- Memory hierarchy

- Why do we have a memory hierarchy
- What is memory hierarchy
- How does memory hierarchy work
- **Locality principle**
  - temporal
  - spatial
- Direct Mapped Cache

reg : fast but small [cap]  
 conflict {  
 MEM: slow but big [cap]

*CPU → cache → MEM*



### Direct-Mapped Cache Example

Start with an empty cache (all blocks are not valid/blank).  
Access the words at the following addresses in sequence.

0	4	8	12	16	12	16	60
001000	000100	001000	001100	010000	001100	010000	111100

0 miss      4 miss      8 miss      12 miss

00	Mem[0]

00	Mem[4]

00	Mem[8]

00	Mem[12]

00	Mem[16]

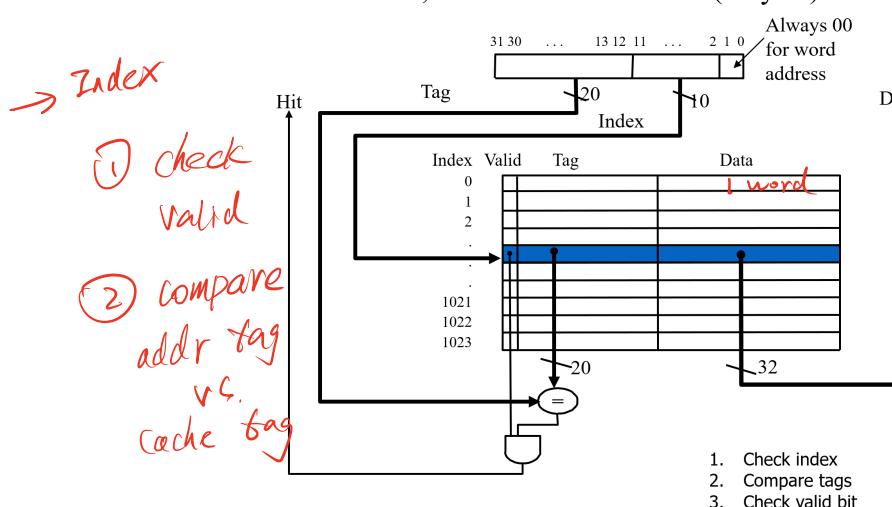
01	Mem[4]

01	Mem[8]

01	Mem[12]

### Direct-Mapped Cache Diagram – Single-word Block

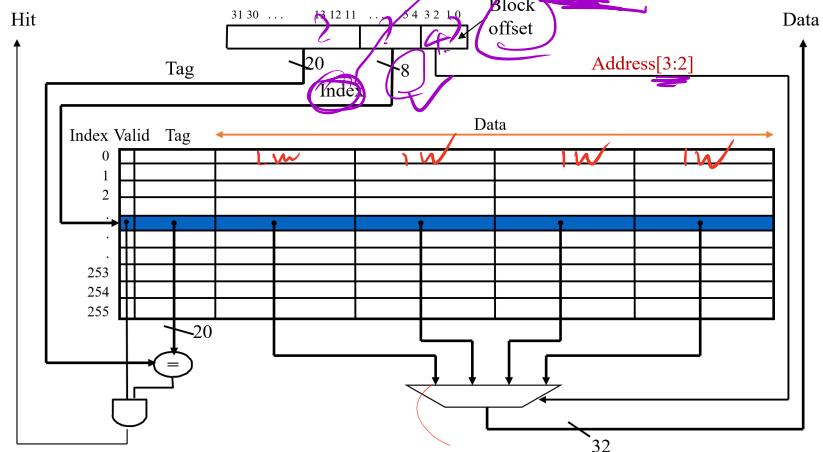
Cache size = 4KiB, block size = one word (4 bytes)



- Multiword Block Direct Mapped Cache

## Direct-Mapped Cache Diagram – Multiword Block

Cache size = 4KiB, block size = four words (16 bytes)



### Another cache of the same size

Let cache block hold two words (8 bytes). Again, start with an empty cache.

0	4	8	12	16	12	16	60
000000	000100	001000	001100	010000	001100	010000	111100

0 miss

4 hit

8 miss

00	Mem[4]	Mem[0]

00	Mem[4]	Mem[0]

00	Mem[4]	Mem[0]
00	Mem[12]	Mem[8]

12 hit

00	Mem[4]	Mem[0]
00	Mem[12]	Mem[8]

01	Mem[4]	Mem[0]
00	Mem[12]	Mem[8]

16 miss

01	Mem[4]	Mem[0]
00	Mem[12]	Mem[8]

01	Mem[20]	Mem[16]
00	Mem[12]	Mem[8]

16 hit

01	Mem[20]	Mem[16]
00	Mem[12]	Mem[8]

60 miss

01	Mem[20]	Mem[16]
00	Mem[12]	Mem[8]

11 60 56

- Handling Write Hits
  - Write-through ①+②
  - Write Back ①
- Handling Miss
  - Write allocate or No write allocate
- Cache Performance

*CPU → Cache → MEM*  
① ? ② ? Data-A

$$\text{MemoryStallCycles} = \frac{\text{Accesses/Program}}{\text{Miss Rate}} \times \text{Miss Penalty}$$

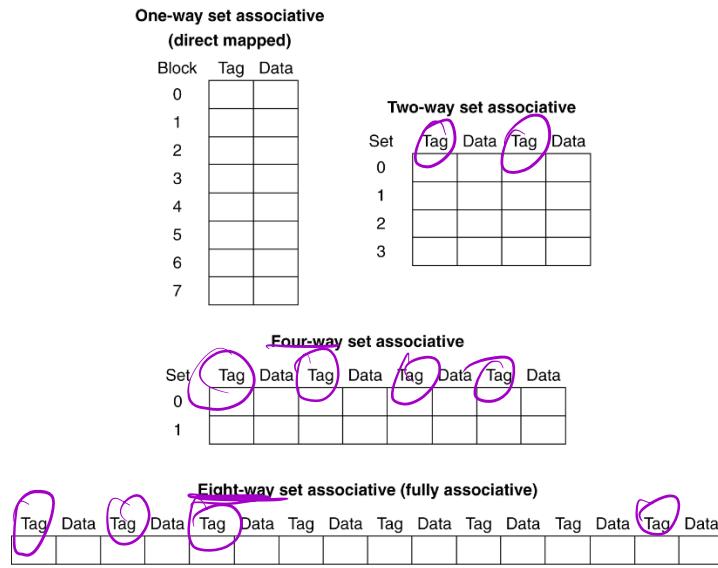
$$\text{MemoryStallCycles} = \text{IC} \times \frac{\text{Misses/Instruction}}{\text{Instruction}} \times \text{Miss Penalty}$$

$$\text{CPU time} = \text{IC} \times (\text{CPI}_{\text{perfect}} + \frac{\text{MemoryStallCycles}}{\text{Instruction}}) \times \text{CC}$$

- Average Memory Access Time (AMAT)

$$\text{AMAT} = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$$

- n-way set associative cache



- Multi-level Cache

